

Sample Size Determination: Posterior Distributions Proximity

Nikita Kiselev^{1*} and Andrey Grabovoy¹

¹Moscow Institute of Physics and Technology, Dolgoprudny, Russia.

*Corresponding author(s). E-mail(s): kiselev.ns@phystech.edu;
Contributing authors: grabovoy.av@phystech.edu;

Abstract

The issue of sample size determination is crucial for constructing an effective machine learning model. However, the existing methods for determining a sufficient sample size are either not strictly proven, or relate to the specific statistical hypothesis about the distribution of model parameters. In this paper we present two approaches based on the proximity of posterior distributions of model parameters on similar subsamples. We show that these two methods are valid for the model with normal posterior distribution of parameters. Computational experiments demonstrate the convergence of the proposed functions as the sample size increases.

Keywords: Sufficient sample size, Posterior distributions proximity, Normal posterior distribution, Linear regression

1 Introduction

The task of supervised machine learning involves selecting a predictive model from a parametric family. This choice is usually based on certain statistical hypotheses, such as maximizing a quality functional. A model that satisfies these statistical hypotheses is called an *adequate* model [1–3].

When planning a computational experiment, it is necessary to estimate the minimum sample size — the number of objects required to build an adequate model. The sample size required to build an adequate predictive model is called *sufficient* [4–6].

This work addresses the issue of determining the sufficient sample size. There are numerous studies dedicated to this topic, with approaches classified into statistical, Bayesian, and heuristic methods.

Some of the early researches on this topic [7, 8] formulate a specific statistical criterion, where the sample size estimation method associated with this criterion guarantees achieving a fixed statistical power with a Type I error not exceeding a specified value. Statistical methods include the Lagrange multipliers test [9], the Likelihood ratio test [10], the Wald statistic [11]. Statistical methods have certain limitations associated with their practical application. They allow for estimating the sample size based on assumptions about the data distribution and information about the agreement of observed values with the assumptions of the null hypothesis.

The Bayesian approach also has a place in this problem. In the work [12] the sufficient sample size is determined based on maximizing the expected utility function. This may explicitly include parameter distribution functions and penalties for increasing the sample size. This work also considers alternative approaches based on constraining a certain quality criterion for estimating model parameters. Among these criteria, the Average Posterior Variance Criterion (APVC), Average Coverage Criterion (ACC), Average Length Criterion (ALC), and Effective Sample Size Criterion (ESC) stand out. These criteria have been further developed in other works, for example, [13] and [14]. Over time, the authors of [15] conducted a theoretical and practical comparison of methods from [7, 8, 12].

Authors like [16], as well as [17], discuss the differences between Bayesian and frequentist approaches in determining sample size. They also propose robust methods for the Bayesian approach and provide illustrative examples for some probabilistic models.

In the paper [18], various methods for estimating sample size in generalized linear models are considered, including statistical, heuristic, and Bayesian methods. Methods such as Lagrange Multiplier Test, Likelihood Ratio Test, Wald Test, Cross-Validation, Bootstrap, Kullback-Leibler Criterion, Average Posterior Variance Criterion, Average Coverage Criterion, Average Length Criterion, and Utility Maximization are analyzed. The authors point out the potential development of combining Bayesian and statistical approaches to estimate sample size for insufficient available sample sizes.

In [19] a method for determining sample size in logistic regression is discussed, using cross-validation and Kullback-Leibler divergence between posterior distributions of model parameters on similar subsamples. Similar subsamples are those that can be obtained from each other by adding, removing, or replacing one object.

In this paper, two approaches based on the distance between the posterior distributions are presented. It is proposed to consider two similar subsamples. The posterior distributions of the model parameters over these subsamples turn out to be close if the sample size is sufficient. It is proposed to use the Kullback-Leibler divergence [19] as a measure of the proximity of distributions, as well as the s-score model comparison function [20]. The novelty of this work lies in proving the correctness of the proposed methods. Correctness is proved in a probabilistic model with a normal posterior distribution of parameters. For the linear regression model, the theorem on the moments of the limit posterior distribution of parameters is proved.

2 Problem statement

An object is defined as a pair (\mathbf{x}, y) , where $\mathbf{x} \in \mathbb{X} \subseteq \mathbb{R}^n$ is the feature vector, and $y \in \mathbb{Y}$ is the target variable. In regression problems $\mathbb{Y} = \mathbb{R}$, and in K -class classification problems $\mathbb{Y} = \{1, \dots, K\}$.

The feature-object matrix for a sample $\mathfrak{D}_m = \{(\mathbf{x}_i, y_i)\}, i \in \mathcal{I} = \{1, \dots, m\}$ of size m is called the matrix $\mathbf{X}_m = [\mathbf{x}_1, \dots, \mathbf{x}_m]^\top \in \mathbb{R}^{m \times n}$.

The target variable vector for a sample $\mathfrak{D}_m = \{(\mathbf{x}_i, y_i)\}, i \in \mathcal{I} = \{1, \dots, m\}$ of size m is denoted by $\mathbf{y}_m = [y_1, \dots, y_m]^\top \in \mathbb{Y}^m$.

A model is a parametric family of functions f , mapping the Cartesian product of the set of feature vector values \mathbb{X} and the set of parameter values \mathbb{W} to the set of target variable values \mathbb{Y} :

$$f : \mathbb{X} \times \mathbb{W} \rightarrow \mathbb{Y}.$$

A probabilistic model is a joint distribution

$$p(y, \mathbf{w}|\mathbf{x}) = p(y|\mathbf{x}, \mathbf{w})p(\mathbf{w}) : \mathbb{Y} \times \mathbb{W} \times \mathbb{X} \rightarrow \mathbb{R}^+,$$

where $\mathbf{w} \in \mathbb{W}$ is the set of model parameters, $p(y|\mathbf{x}, \mathbf{w})$ specifies the likelihood of an object, and $p(\mathbf{w})$ represents the prior distribution of parameters.

The task is to determine the sufficient sample size m^* . Let a criterion T be given. E.g. it can be constructed based on heuristics regarding the behaviour of model parameters.

Definition 1. *The sample size m^* is called **sufficient** according to the criterion T , if T holds for all $k \geq m^*$.*

3 Proposed sample size determination methods

In [19], it is suggested to use the Kullback-Leibler divergence to estimate a sufficient sample size in a binary classification problem. The idea is based on the fact that if two subsamples differ from each other by one object, then the posterior distributions obtained from them should be close. This proximity is determined by the Kullback-Leibler divergence.

In this paper, the question of the correctness of this approach is considered. The method is studied in an arbitrary probabilistic model. As a measure of proximity, it is proposed to use not only the Kullback-Leibler divergence, but also the s-score similarity function from [20].

Consider two subsamples $\mathfrak{D}^1 \subseteq \mathfrak{D}_m$ and $\mathfrak{D}^2 \subseteq \mathfrak{D}_m$. Let $\mathcal{I}_1 \subseteq \mathcal{I} = \{1, \dots, m\}$ and $\mathcal{I}_2 \subseteq \mathcal{I} = \{1, \dots, m\}$ — corresponding to them subsets of indexes.

Definition 2. *Subsamples \mathfrak{D}^1 and \mathfrak{D}^2 are called **similar** if \mathcal{I}_2 can be obtained from \mathcal{I}_1 by deleting, replacing or adding one element, that is*

$$|\mathcal{I}_1 \triangle \mathcal{I}_2| = |(\mathcal{I}_1 \setminus \mathcal{I}_2) \cup (\mathcal{I}_2 \setminus \mathcal{I}_1)| = 1.$$

Consider two similar subsamples $\mathfrak{D}_k = (\mathbf{X}_k, \mathbf{y}_k)$ and $\mathfrak{D}_{k+1} = (\mathbf{X}_{k+1}, \mathbf{y}_{k+1})$ of sizes k and $k+1$, respectively. This means that the larger one is obtained by adding

one element to the smaller one. Let's find the posterior distribution of the model parameters over these subsamples:

$$p_j(\mathbf{w}) = p(\mathbf{w}|\mathfrak{D}_j) = \frac{p(\mathfrak{D}_j|\mathbf{w})p(\mathbf{w})}{p(\mathfrak{D}_j)} \propto p(\mathfrak{D}_j|\mathbf{w})p(\mathbf{w}), \quad j = k, k+1.$$

Definition 3. Let's fix some positive number $\varepsilon > 0$. The sample size m^* is called **KL-sufficient** if for all $k \geq m^*$

$$KL(k) = D_{KL}(p_k \| p_{k+1}) = \int p_k(\mathbf{w}) \log \frac{p_k(\mathbf{w})}{p_{k+1}(\mathbf{w})} d\mathbf{w} \leq \varepsilon.$$

For a pair of normal distributions, the Kullback-Leibler divergence has a fairly simple form. Assume that the posterior distribution is normal, that is, $p_k(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_k, \mathbf{\Sigma}_k)$. Guided by the heuristic that the convergence of the moments of such a distribution should entail the proximity of posterior distributions on similar subsamples, the following statement can be formulated.

Theorem 1. Let $\|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2 \rightarrow 0$ and $\|\mathbf{\Sigma}_{k+1} - \mathbf{\Sigma}_k\|_F \rightarrow 0$ as $k \rightarrow \infty$. Then, in a model with a normal posterior distribution of parameters, the definition of a KL-sufficient sample size is correct. Namely, for any $\varepsilon > 0$, there is such a m^* that for all $k \geq m^*$ $KL(k) \leq \varepsilon$ is satisfied.

In this paper, it is proposed to use the s-score similarity function from [20] as a measure of proximity of distributions:

$$\text{s-score}(g_1, g_2) = \frac{\int_{\mathbf{w}} g_1(\mathbf{w}) g_2(\mathbf{w}) d\mathbf{w}}{\max_{\mathbf{b}} \int_{\mathbf{w}} g_1(\mathbf{w} - \mathbf{b}) g_2(\mathbf{w}) d\mathbf{w}}.$$

Definition 4. Let's fix some positive number $\varepsilon > 0$. The sample size m^* is called **S-sufficient** if for all $k \geq m^*$

$$S(k) = \text{s-score}(p_k, p_{k+1}) \geq 1 - \varepsilon.$$

As in the case of a KL-sufficient sample size, in a model with a normal posterior distribution, it is possible to write an expression for the criterion used. Thus, one more statement can be formulated.

Theorem 2. Let $\|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2 \rightarrow 0$ as $k \rightarrow \infty$. Then, in a model with a normal posterior distribution of parameters, the definition of an S-sufficient sample size is correct. Namely, for any $\varepsilon > 0$, there is such a m^* that for all $k \geq m^*$ $S(k) \geq 1 - \varepsilon$ is satisfied.

Let the linear regression model have a normal prior distribution of parameters. By the conjugacy property of the prior distribution and likelihood, the posterior distribution is also normal. Thus, we come to one of the simplest examples of a model for which the theorems presented above are valid. In fact, simpler statements can be formulated for linear regression.

Theorem 3. Let the sets of values of the features and the target variable be bounded, that is, $\exists M \in \mathbb{R} : \|\mathbf{x}\|_2 \leq M$ and $|y| \leq M$. If $\lambda_{\min}(\mathbf{X}_k^T \mathbf{X}_k) = \omega(\sqrt{k})$ for $k \rightarrow$

∞ , then in a linear regression model with a normal prior distribution of parameters $\|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2 \rightarrow 0$ and $\|\boldsymbol{\Sigma}_{k+1} - \boldsymbol{\Sigma}_k\|_F \rightarrow 0$ as $k \rightarrow \infty$.

4 Computational experiment

This section provides an empirical study of the proposed methods. Experiments were conducted on synthetic data and dataset Liver Disorders from [21].

Synthetic data is generated from a linear regression model. The number of objects is 500, the number of features is 10. One object is sequentially removed from the given sample until the number of objects in the subsample is equal to the number of features. For each sample size k we calculate the minimum eigenvalue of the matrix $\mathbf{X}_k^\top \mathbf{X}_k$. Also, the values of $KL(k)$ and $S(k)$ are calculated. This process is repeated $B = 100$ times.

Regression dataset Liver Disorders has 345 objects and 5 features. We also sequentially remove objects from sample one by one. Minimum eigenvalue and function values are calculated. This process is repeated $B = 1000$ times.

Fig. 1 shows the asymptotic behavior of the minimum eigenvalue of the matrix $\mathbf{X}_k^\top \mathbf{X}_k$. We see that when the sample size tends to infinity, the minimum eigenvalue also tends to infinity. Meanwhile, as is necessary for the Theorem 3, the graph is higher than \sqrt{k} .




Fig. 1 Minimum eigenvalue vs available sample size

In the Fig. 2, we can observe the obtained dependencies between the available sample size k and the proposed functions $KL(k)$ and $S(k)$ for the synthetic regression dataset. At the same time, in the Fig. 3, we see the same plots for the Liver Disorders dataset. It can be seen that in both cases, the value of $KL(k)$ approaches zero as the sample size increases, and $S(k)$ tends towards one. These empirical results confirm the theoretical ones obtained earlier.




Fig. 2 Synthetic regression dataset




Fig. 3 Liver Disorders dataset

For the definitions of KL-sufficiency and S-sufficiency, there is a hyperparameter ε , which corresponds to the threshold for a sufficient sample size m^* . In order to study the dependence between them, we introduce Fig. 4, which shows what sample sizes can be chosen to provide a certain level of confidence.

To compare the performance of the proposed method on different datasets, samples have been chosen from the open repository [21]. The detailed information about each dataset, the number of observations and number of features, are provided in Table 1. For demonstration purposes, a value of the hyperparameter ε has been selected, at which the value of the target function, either $KL(k)$ or $S(k)$, decreases by half. The corresponding results are in Table 1. Omissions mean that the initial sample size is not sufficient.

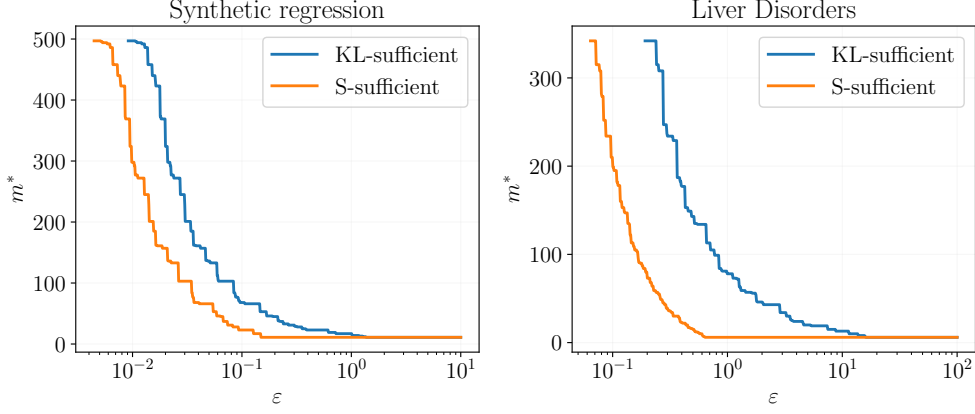


Fig. 4 Sufficient sample size vs threshold

Table 1 Datasets with regression task (omissions mean that the initial sample size is not sufficient)

Dataset name	Objects m	Features n	KL	S
Abalone	4177	8	3751	3751
Auto MPG	392	8	55	—
Automobile	159	25	152	—
Liver Disorders	345	6	282	—
Servo	167	4	138	160
Forest fires	517	12	487	—
Wine Quality	6497	12	—	—
Energy Efficiency	768	9	—	—
Student Performance	649	32	—	—
Facebook Metrics	495	18	379	475
Real Estate Valuation	414	7	—	—
Heart Failure Clinical Records	299	12	258	293
Bone marrow transplant: children	142	36	109	—

5 Conclusion

Approaches to determining a sufficient sample size based on the proximity of posterior distributions of model parameters on similar subsamples are proposed. The correctness of the proposed approaches is proved under certain restrictions on the model used. The theorem on the moments of the limit posterior distribution of parameters in a linear regression model is proved. The conducted computational experiment makes it possible to analyze the properties of the proposed methods and their effectiveness.

Appendix A Proof of Theorem 1

Proof. The Kullback-Leibler divergence for a pair of normal posterior distributions has the form

$$D_{\text{KL}}(p_k \| p_{k+1}) = \frac{1}{2} \left(\text{tr}(\Sigma_{k+1}^{-1} \Sigma_k) + (\mathbf{m}_{k+1} - \mathbf{m}_k)^\top \Sigma_{k+1}^{-1} (\mathbf{m}_{k+1} - \mathbf{m}_k) - n + \log \left(\frac{\det \Sigma_{k+1}}{\det \Sigma_k} \right) \right).$$

Let's express Σ_{k+1} as $\Sigma_{k+1} = \Sigma_k + \Delta \Sigma$. Let's consider each term separately.

$$\text{tr}(\Sigma_{k+1}^{-1} \Sigma_k) = \text{tr}((\Sigma_k + \Delta \Sigma)^{-1} \Sigma_k) \rightarrow \text{tr} \mathbf{I}_n = n \text{ as } \|\Delta \Sigma\|_F \rightarrow 0,$$

$$|(\mathbf{m}_{k+1} - \mathbf{m}_k)^\top \Sigma_{k+1}^{-1} (\mathbf{m}_{k+1} - \mathbf{m}_k)| \leq \|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2^2 \|\Sigma_{k+1}^{-1}\|_2 \rightarrow 0 \text{ as } \|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2 \rightarrow 0,$$

$$\log \left(\frac{\det \Sigma_{k+1}}{\det \Sigma_k} \right) = \log \left(\frac{\det(\Sigma_k + \Delta \Sigma)}{\det \Sigma_k} \right) \rightarrow \log \det \mathbf{I}_n = \log 1 = 0 \text{ as } \|\Delta \Sigma\|_F \rightarrow 0,$$

from where we have the required. \square

Appendix B Proof of Theorem 2

Proof. Let's use the s-score expression for a pair of normal posterior distributions from [20]:

$$\text{s-score}(p_k, p_{k+1}) = \exp \left(-\frac{1}{2} (\mathbf{m}_{k+1} - \mathbf{m}_k)^\top (\Sigma_k + \Sigma_{k+1})^{-1} (\mathbf{m}_{k+1} - \mathbf{m}_k) \right).$$

Because

$$\left| (\mathbf{m}_{k+1} - \mathbf{m}_k)^\top (\Sigma_k + \Sigma_{k+1})^{-1} (\mathbf{m}_{k+1} - \mathbf{m}_k) \right| \leq \|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2^2 \|(\Sigma_k + \Sigma_{k+1})^{-1}\|_2 \rightarrow 0$$

if $\|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2 \rightarrow 0$, then the value of the quadratic form inside the exponent tends to zero. Therefore, $\text{s-score}(p_k, p_{k+1}) \rightarrow 1$ as $\|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2 \rightarrow 0$. \square

Appendix C Proof of Theorem 3

Proof. Let be a normal prior distribution of parameters $p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \alpha^{-1} \mathbf{I})$. In a linear regression model, likelihood is normal, namely

$$p(\mathbf{y} | \mathbf{X}, \mathbf{w}) = \mathcal{N}(\mathbf{y} | \mathbf{X}\mathbf{w}, \sigma^2 \mathbf{I}) = (2\pi\sigma^2)^{-m/2} \exp \left(-\frac{1}{2\sigma^2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 \right).$$

Using the conjugacy of the prior distribution and likelihood, it is easy to find the parameters of the posterior distribution:

$$p(\mathbf{w} | \mathbf{X}, \mathbf{y}) = \mathcal{N}(\mathbf{w} | \mathbf{m}, \Sigma),$$

where

$$\boldsymbol{\Sigma} = \left(\alpha \mathbf{I} + \frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X} \right)^{-1}, \quad \mathbf{m} = (\mathbf{X}^\top \mathbf{X} + \alpha \sigma^2 \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}.$$

Consider the expression $\|\boldsymbol{\Sigma}_{k+1} - \boldsymbol{\Sigma}_k\|_2$ norms of difference of covariance matrices for subsamples of size k and $k+1$. Let's introduce the notation $\mathbf{A}_k = \frac{1}{\sigma^2} \mathbf{X}_k^\top \mathbf{X}_k$. Given the formulas above, we have

$$\begin{aligned} \|\boldsymbol{\Sigma}_{k+1} - \boldsymbol{\Sigma}_k\|_2 &= \left\| (\alpha \mathbf{I} + \mathbf{A}_{k+1})^{-1} - (\alpha \mathbf{I} + \mathbf{A}_k)^{-1} \right\|_2 = \\ &= \left\| (\alpha \mathbf{I} + \mathbf{A}_{k+1})^{-1} (\mathbf{A}_{k+1} - \mathbf{A}_k) (\alpha \mathbf{I} + \mathbf{A}_k)^{-1} \right\|_2 \leq \end{aligned}$$

Let's use the submultiplicativity of the spectral matrix norm.

$$\leq \left\| (\alpha \mathbf{I} + \mathbf{A}_{k+1})^{-1} \right\|_2 \left\| (\alpha \mathbf{I} + \mathbf{A}_k)^{-1} \right\|_2 \|\mathbf{A}_{k+1} - \mathbf{A}_k\|_2 =$$

Now let's use the expression of the spectral matrix norm in terms of the maximum eigenvalue.

$$\begin{aligned} &= \frac{1}{\lambda_{\min}(\alpha \mathbf{I} + \mathbf{A}_{k+1})} \frac{1}{\lambda_{\min}(\alpha \mathbf{I} + \mathbf{A}_k)} \|\mathbf{A}_{k+1} - \mathbf{A}_k\|_2 \leq \\ &\leq \frac{1}{\lambda_{\min}(\mathbf{A}_{k+1})} \frac{1}{\lambda_{\min}(\mathbf{A}_k)} \|\mathbf{A}_{k+1} - \mathbf{A}_k\|_2 = \\ &= \sigma^2 \frac{1}{\lambda_{\min}(\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1})} \frac{1}{\lambda_{\min}(\mathbf{X}_k^\top \mathbf{X}_k)} \|\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1} - \mathbf{X}_k^\top \mathbf{X}_k\|_2. \end{aligned}$$

Further, since by the condition $\|\mathbf{x}\|_2 \leq M$, then

$$\|\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1} - \mathbf{X}_k^\top \mathbf{X}_k\|_2 = \left\| \sum_{i=1}^{k+1} \mathbf{x}_i \mathbf{x}_i^\top - \sum_{i=1}^k \mathbf{x}_i \mathbf{x}_i^\top \right\|_2 = \|\mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top\|_2 = \lambda_{\max}(\mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top) =$$

A matrix of unit rank has a single nonzero eigenvalue.

$$= \mathbf{x}_{k+1}^\top \mathbf{x}_{k+1} = \|\mathbf{x}_{k+1}\|_2^2 \leq M^2.$$

By condition $\lambda_{\min}(\mathbf{X}_k^\top \mathbf{X}_k) = \omega(\sqrt{k})$, then $\|\boldsymbol{\Sigma}_{k+1} - \boldsymbol{\Sigma}_k\|_2 = o(k^{-1})$ as $k \rightarrow \infty$. Next, we will use the equivalence of matrix norms, namely

$$\|\boldsymbol{\Sigma}_{k+1} - \boldsymbol{\Sigma}_k\|_F \leq \sqrt{k} \|\boldsymbol{\Sigma}_{k+1} - \boldsymbol{\Sigma}_k\|_2 = o(k^{-1/2}) \text{ as } k \rightarrow \infty,$$

which was exactly what needed to be proved. Now let's estimate the norm of the difference in mathematical expectations.

$$\|\mathbf{m}_{k+1} - \mathbf{m}_k\|_2 = \left\| (\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1} + \alpha \sigma^2 \mathbf{I})^{-1} \mathbf{X}_{k+1}^\top \mathbf{y}_{k+1} - (\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I})^{-1} \mathbf{X}_k^\top \mathbf{y}_k \right\|_2 =$$

Consider that $\mathbf{X}_{k+1}^\top = [\mathbf{X}_k^\top, \mathbf{x}_{k+1}]$ and $\mathbf{y}_{k+1} = [\mathbf{y}_k, y_{k+1}]^\top$, then $\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1} = \mathbf{X}_k^\top \mathbf{X}_k + \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top$ and $\mathbf{X}_{k+1}^\top \mathbf{y}_{k+1} = \mathbf{X}_k^\top \mathbf{y}_k + \mathbf{x}_{k+1} y_{k+1}$.

$$= \left\| \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} + \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)^{-1} (\mathbf{X}_k^\top \mathbf{y}_k + \mathbf{x}_{k+1} y_{k+1}) - \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{X}_k^\top \mathbf{y}_k \right\|_2 =$$

Let's take out the multiplier in the first term:

$$\left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} + \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)^{-1} = \left(\mathbf{I} + \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)^{-1} \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1}.$$

Next, we will take out the common multiplier for both terms.

$$= \left\| \left[\left(\mathbf{I} + \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)^{-1} - \mathbf{I} \right] \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{X}_k^\top \mathbf{y}_k + \left(\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1} + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} y_{k+1} \right\|_2 =$$

Let's use the triangle inequality, as well as the consistency and submultiplicativity property of the spectral norm.

$$\leq \left\| \left(\mathbf{I} + \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)^{-1} - \mathbf{I} \right\|_2 \left\| \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \right\|_2 \left\| \mathbf{X}_k^\top \mathbf{y}_k \right\|_2 + \left\| \left(\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1} + \alpha \sigma^2 \mathbf{I} \right)^{-1} \right\|_2 \left\| \mathbf{x}_{k+1} y_{k+1} \right\|_2$$

Let's evaluate each term separately. In the first multiplier of the first term, we apply the formula for the difference of inverse matrices, as we did with covariance matrices.

$$\begin{aligned} & \left\| \left(\mathbf{I} + \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)^{-1} - \mathbf{I} \right\|_2 \leq \\ & \leq \left\| \left(\mathbf{I} + \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)^{-1} \right\|_2 \cdot \left\| \mathbf{I} \right\|_2 \cdot \left\| \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right\|_2 \leq \end{aligned}$$

Again, we use submultiplicativity, as well as an expression for the norm of a matrix of unit rank.

$$\begin{aligned} & \leq \frac{1}{\lambda_{\min} \left(\mathbf{I} + \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)} \frac{\left\| \mathbf{x}_{k+1} \right\|_2^2}{\lambda_{\min} \left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)} \leq \\ & \leq \frac{1}{1 + \lambda_{\min} \left(\left(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I} \right)^{-1} \mathbf{x}_{k+1} \mathbf{x}_{k+1}^\top \right)} \frac{M^2}{\lambda_{\min} \left(\mathbf{X}_k^\top \mathbf{X}_k \right)} \leq \end{aligned}$$

The minimum eigenvalue of the product of matrices is estimated by the product of their minimum eigenvalues. In addition, the minimum eigenvalue of the matrix of unit rank $\mathbf{x}_{k+1}\mathbf{x}_{k+1}^\top$ is zero.

$$\leq \frac{1}{1 + \lambda_{\max}(\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I})} \frac{M^2}{\lambda_{\min}(\mathbf{x}_{k+1}\mathbf{x}_{k+1}^\top)} \frac{M^2}{\lambda_{\min}(\mathbf{X}_k^\top \mathbf{X}_k)} = \frac{M^2}{\lambda_{\min}(\mathbf{X}_k^\top \mathbf{X}_k)}.$$

The second and third multipliers of the first term are evaluated as follows.

$$\left\| (\mathbf{X}_k^\top \mathbf{X}_k + \alpha \sigma^2 \mathbf{I})^{-1} \right\|_2 \left\| \mathbf{X}_k^\top \mathbf{y}_k \right\|_2 \leq \frac{\left\| \mathbf{X}_k^\top \mathbf{y}_k \right\|_2}{\lambda_{\min}(\mathbf{X}_k^\top \mathbf{X}_k)} = \frac{\left\| \sum_{i=1}^k \mathbf{x}_i y_i \right\|_2}{\lambda_{\min}(\mathbf{X}_k^\top \mathbf{X}_k)} \leq \frac{k M^2}{\lambda_{\min}(\mathbf{X}_k^\top \mathbf{X}_k)}$$

Finally, let's evaluate the second term.

$$\left\| (\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1} + \alpha \sigma^2 \mathbf{I})^{-1} \right\|_2 \left\| \mathbf{x}_{k+1} y_{k+1} \right\|_2 \leq \frac{M^2}{\lambda_{\min}(\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1})}$$

In total, we have the following estimate.

$$\left\| \mathbf{m}_{k+1} - \mathbf{m}_k \right\|_2 \leq \frac{k M^3}{\lambda_{\min}^2(\mathbf{X}_k^\top \mathbf{X}_k)} + \frac{M^2}{\lambda_{\min}(\mathbf{X}_{k+1}^\top \mathbf{X}_{k+1})} = k \cdot o(k^{-1}) + o(k^{-1/2}) = o(1) \text{ as } k \rightarrow \infty$$

Thus, we obtained the required convergence. \square

References

- [1] Bies, R.R., Muldoon, M.F., Pollock, B.G., Manuck, S., Smith, G., Sale, M.E.: A genetic algorithm-based, hybrid machine learning approach to model selection. *Journal of pharmacokinetics and pharmacodynamics* **33**(2), 195 (2006)
- [2] Cawley, G.C., Talbot, N.L.: On over-fitting in model selection and subsequent selection bias in performance evaluation. *The Journal of Machine Learning Research* **11**, 2079–2107 (2010)
- [3] Raschka, S.: Model evaluation, model selection, and algorithm selection in machine learning. *arXiv preprint arXiv:1811.12808* (2018)
- [4] Byrd, R.H., Chin, G.M., Nocedal, J., Wu, Y.: Sample size selection in optimization methods for machine learning. *Mathematical programming* **134**(1), 127–155 (2012)
- [5] Figueroa, R.L., Zeng-Treitler, Q., Kandula, S., Ngo, L.H.: Predicting sample size required for classification performance. *BMC medical informatics and decision making* **12**, 1–10 (2012)

- [6] Balki, I., Amirabadi, A., Levman, J., Martel, A.L., Emersic, Z., Meden, B., Garcia-Pedrero, A., Ramirez, S.C., Kong, D., Moody, A.R., *et al.*: Sample-size determination methodologies for machine learning in medical imaging research: a systematic review. *Canadian Association of Radiologists Journal* **70**(4), 344–353 (2019)
- [7] Adcock, C.J.: A bayesian approach to calculating sample sizes. *The Statistician* **37**(4/5), 433 (1988) <https://doi.org/10.2307/2348770>
- [8] Joseph, L., Wolfson, D.B., Berger, R.D.: Sample size calculations for binomial proportions via highest posterior density intervals. *Journal of the Royal Statistical Society. Series D (The Statistician)* **44**(2), 143–154 (1995). Accessed 2023-12-12
- [9] Self, S.G., Mauritsen, R.H.: Power/sample size calculations for generalized linear models. *Biometrics*, 79–86 (1988)
- [10] Shieh, G.: On power and sample size calculations for likelihood ratio tests in generalized linear models. *Biometrics* **56**(4), 1192–1196 (2000)
- [11] Shieh, G.: On power and sample size calculations for wald tests in generalized linear models. *Journal of Statistical Planning and Inference* **128**(1), 43–59 (2005)
- [12] Lindley, D.V.: The choice of sample size. *Journal of the Royal Statistical Society: Series D (The Statistician)* **46**(2), 129–138 (1997) <https://doi.org/10.1111/1467-9884.00068>
- [13] Pham-Gia, T.: On bayesian analysis, bayesian decision theory and the sample size problem. *Journal of the Royal Statistical Society: Series D (The Statistician)* **46**(2), 139–144 (1997) <https://doi.org/10.1111/1467-9884.00069>
- [14] Gelfand, A.E., Wang, F.: A simulation-based approach to bayesian sample size determination for performance under a given model and for separating models. *Statistical Science* **17**(2) (2002) <https://doi.org/10.1214/ss/1030550861>
- [15] Cao, J., Lee, J.J., Alber, S.: Comparison of bayesian sample size criteria: Acc, alc, and woc. *Journal of Statistical Planning and Inference* **139**(12), 4111–4122 (2009) <https://doi.org/10.1016/j.jspi.2009.05.041>
- [16] Brutti, P., De Santis, F., Gubbiotti, S.: Bayesian-frequentist sample size determination: a game of two priors. *METRON* **72**(2), 133–151 (2014) <https://doi.org/10.1007/s40300-014-0043-2>
- [17] Pezeshk, H., Nematollahi, N., Maroufy, V., Gittins, J.: The choice of sample size: a mixed bayesian / frequentist approach. *Statistical Methods in Medical Research* **18**(2), 183–194 (2008) <https://doi.org/10.1177/0962280208089298>
- [18] Grabovoy, A.V., Gadaev, T.T., Motrenko, A.P., Strijov, V.V.: Numerical

methods of sufficient sample size estimation for generalised linear models. Lobachevskii Journal of Mathematics **43**(9), 2453–2462 (2022) <https://doi.org/10.1134/s1995080222120125>

- [19] Motrenko, A., Strijov, V., Weber, G.-W.: Sample size determination for logistic regression. Journal of Computational and Applied Mathematics **255**, 743–752 (2014) <https://doi.org/10.1016/j.cam.2013.06.031>
- [20] Aduenko, A.: Selection of multimodels in classification tasks. PhD thesis, MIPT (2017). https://www.frccsc.ru/diss-council/00207305/diss/list/aduenko_aa
- [21] Dua, D., Graff, C.: UCI Machine Learning Repository (2017). <http://archive.ics.uci.edu/ml>